

# PARALLEL MEMORY ARCHITECTURE

Mahdi Nazm Bojnordi



Assistant Professor

School of Computing

University of Utah

# Chip Multiprocessors

- Can be viewed as a simple SMP on single chip
- CPUs are now called cores
  - ▣ One thread per core
- Shared higher level caches
  - ▣ Typically the last level
  - ▣ Lower latency
  - ▣ Improved bandwidth
- Not necessarily homogenous cores!



*Intel Nehalem (Core i7)*

# Efficiency of Chip Multiprocessing

- **Ideally**,  $n$  cores provide  $n\times$  performance
- Example: design an ideal dual-processor
  - ▣ **Goal**: provide the same performance as uniprocessor

	Uniprocessor	Dual-processor
Frequency	1	?
Voltage	1	?
Execution Time	1	1
Dynamic Power	1	?
Dynamic Energy	1	?
Energy Efficiency	1	?

# Efficiency of Chip Multiprocessing

- **Ideally**,  $n$  cores provide  $n\times$  performance
- Example: design an ideal dual-processor
  - ▣ **Goal**: provide the same performance as uniprocessor

$$f \propto V \text{ \& } P \propto V^3 \rightarrow V_{dual} = 0.5V_{uni} \rightarrow P_{dual} = 2 \times 0.125P_{uni}$$

	Uniprocessor	Dual-processor
Frequency	1	0.5
Voltage	1	0.5
Execution Time	1	1
Dynamic Power	1	$2 \times 0.125$
Dynamic Energy	1	$2 \times 0.125$
Energy Efficiency	1	4

# Challenges

# Example Code I

- A sequential application runs as a single thread


**Kernel Function:**

```
void kern (int start, int end) {  
    int i;  
    for(i=start; i<=end; ++i) {  
        A[i] = A[i] * A[i] + 5;  
    }  
}
```


**Single Thread**

```
main() {  
    ...  
    kern (1, n);  
    ...  
}
```

**Memory**



**Processor**



# Example Code I

- Two threads operating on separate partitions


**Kernel Function:**

```
void kern (int start, int end) {  
    int i;  
    for(i=start; i<=end; ++i) {  
        A[i] = A[i] * A[i] + 5;  
    }  
}
```

**Thread 0**

```
main() {  
    ...  
    kern (1, n/2);  
    ...  
}
```

**Memory**



**Processor**

**Processor**

**Thread 1**


```
kern (n/2+1, n);
```

# Performance of Parallel Processing

- Recall: Amdahl's law for theoretical speedup
  - ▣ Overall speedup is limited to the fraction of the program that can be executed in parallel

$$speedup = \frac{1}{f + \frac{1-f}{n}} \quad f: \text{sequential fraction}$$

**Speedup vs. Sequential Fraction**



— 10% — 20% — 40% — 60% — 90%



# Example Code II

- A single location is updated every time


**Kernel Function:**

```
void kern (int start, int end) {  
    int i;  
    for(i=start; i<=end; ++i) {  
        sum = sum + A[i];  
    }  
}
```



**Thread 0**

```
main() {  
    ...  
    kern (1, n);  
    ...  
}
```

**Memory**



**Processor**



# Example Code II

- A single location is updated every time


**Kernel Function:**

```
void kern (int start, int end) {  
    int i;  
    for(i=start; i<=end; ++i) {  
        sum = sum + A[i];  
    }  
}
```


**Thread 0**

```
main() {  
    ...  
    kern (1, n);  
    ...  
}
```

**Memory**



**Processor**



# Example Code II

- Two threads operating on separate partitions


Kernel Function:

```
void kern (int start, int end) {  
    int i;  
    for(i=start; i<=end; ++i) {  
        sum = sum + A[i];  
    }  
}
```

Thread 0

```
main() {  
    ...  
    kern (1, n/2);  
    ...  
}
```

Memory



Processor

Processor

Thread 1


```
kern (n/2+1, n);
```

# Communication in Multiprocessors

- How multiple processor cores communicate?


## Shared Memory

- Multiple threads employ shared memory
- Easy for programmers (loads and stores)



## Message Passing

- Explicit communication through interconnection network
- Simple hardware




# Shared Memory Architectures

## Uniform Memory Access

- Equal latency for all processors
- Simple software control


### Example UMA



## Non-Uniform Memory Access

- Access latency is proportional to proximity
  - ▣ Fast local accesses

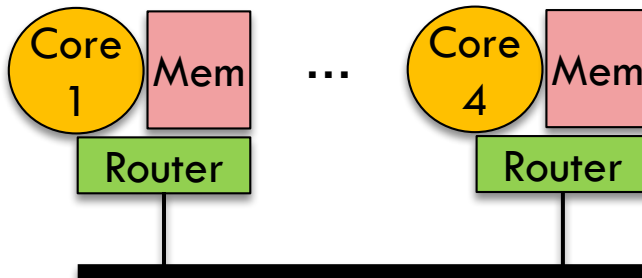
### Example NUMA



# Network Topologies


## Shared Network

- ❑ Low latency
- ❑ Low bandwidth
- ❑ Simple control
  - ▣ e.g., bus



## Point to Point Network

- ❑ High latency
- ❑ High bandwidth
- ❑ Complex control
  - ▣ e.g., mesh, ring




# Challenges in Shared Memories

- Correctness of an application is influenced by
  - ▣ Memory consistency
    - All memory instructions appear to execute in the **program order**
    - Known to the programmer
  - ▣ Cache coherence
    - All the processors see the **same data** for a particular memory address as they should have if there were no caches in the system
    - Invisible to the programmer

# Cache Coherence Problem


- ❑ Multiple copies of each cache block
  - ▣ In main memory and caches
- ❑ Multiple copies can get inconsistent when writes happen
  - ▣ Solution: propagate writes from one core to others






# Scenario 1: Loading From Memory

- ❑ Variable A initially has value 0
- ❑ P1 stores value 1 into A
- ❑ P2 loads A from memory and sees old value 0



# Scenario 2: Loading From Cache

- ❑ P1 and P2 both have variable A (value 0) in their caches
- ❑ P1 stores value 1 into A
- ❑ P2 loads A from its cache and sees old value





# Cache Coherence

- The key operation is **update/invalidate** sent to all or a subset of the cores
  - ▣ Software based management
    - Flush: write all of the dirty blocks to memory
    - Invalidate: make all of the cache blocks invalid
  - ▣ Hardware based management
    - Update or invalidate other copies on every write
    - Send data to everyone, or only the ones who have a copy
- Invalidation based protocol is better. **Why?**


# Snoopy Protocol

- ❑ Relying on a broadcast infrastructure among caches
  - ▣ For example shared bus
- ❑ Every cache monitors (**snoop**) the traffic on the shared media to keep the states of the cache block up to date




# Simple Snooping Protocol

- Relies on write-through, write no-allocate cache
- Multiple readers are allowed
  - ▣ Writes invalidate replicas
- Employs a simple state machine for each cache unit



# Simple Snooping State Machine

- Every node updates its one-bit valid flag using a simple finite state machine (FSM)
- Processor actions
  - ▣ Load, Store, Evict
- Bus traffic
  - ▣ BusRd, BusWr



→ Transaction by local actions  
-----> Transaction by bus traffic

# Snooping with Writeback Policy



- **Problem:** writes are not propagated to memory until eviction
  - ▣ Cache data maybe different from main memory
- **Solution:** identify the **owner** of the most recently updated replica
  - ▣ Every data may have only one owner at any time
  - ▣ Only the owner can update the replica
  - ▣ Multiple readers can share the data
    - No one can write without gaining ownership first

# Modified-Shared-Invalid Protocol

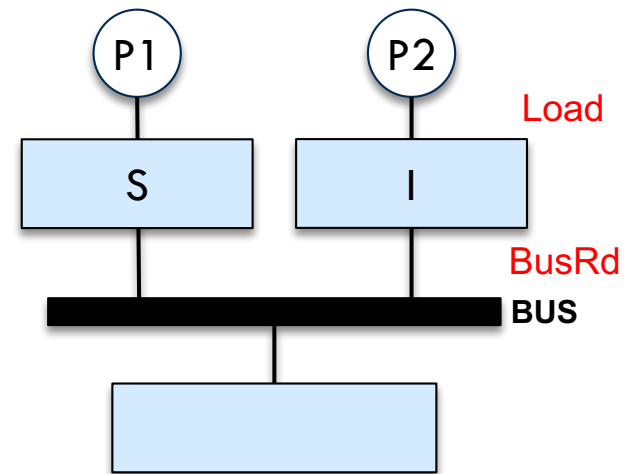

- Every cache block transitions among three states
  - ▣ **Invalid**: no replica in the cache
  - ▣ **Shared**: a read-only copy in the cache
    - Multiple units may have the same copy
  - ▣ **Modified**: a writable copy of the data in the cache
    - The replica has been updated
    - The cache has the only valid copy of the data block
- Processor actions
  - ▣ Load, store, evict
- Bus messages
  - ▣ BusRd, BusRdX, BusInv, BusWB, BusReply




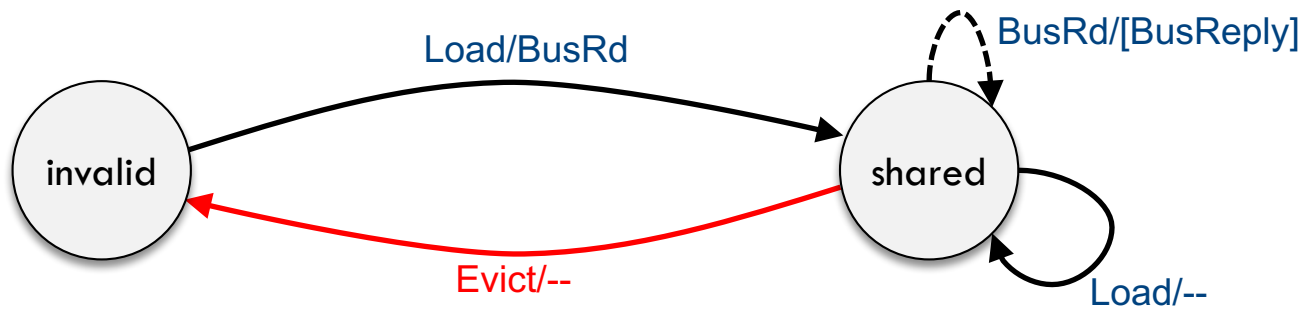
# MSI Example



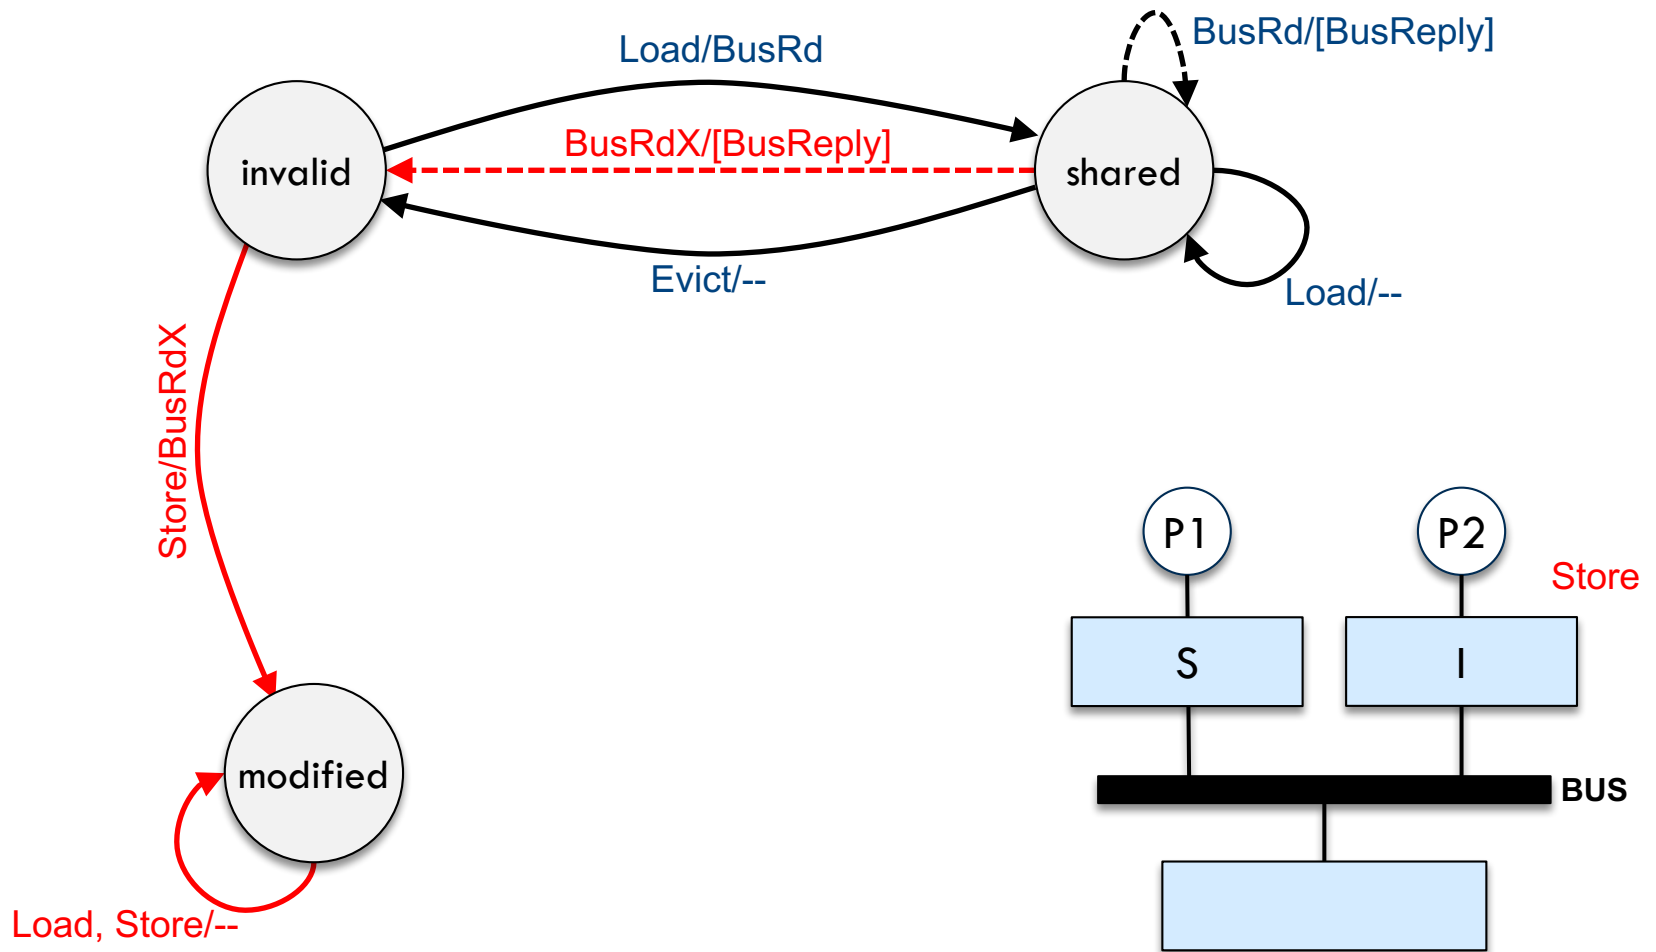
# MSI Example




# MSI Example



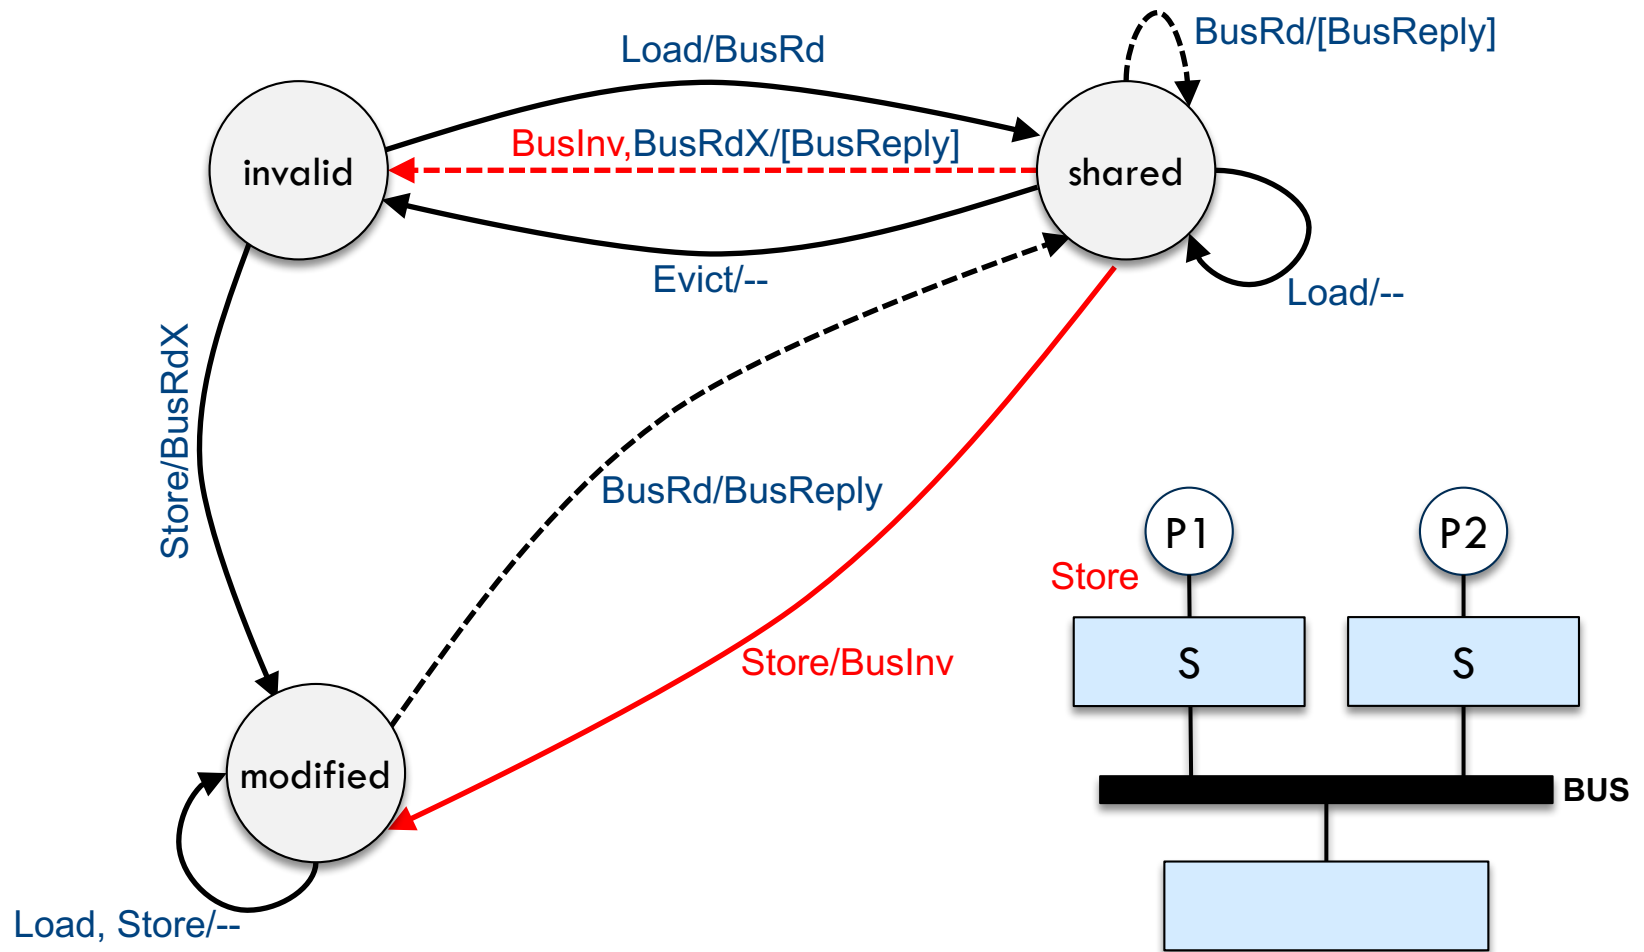
# MSI Example




# MSI Example



# MSI Example



# MSI Example



# MSI Example

